


Will alignment techniques scale to RL?





Daniel Tan

Alignment research, 2024–2026

MATS · W/ OWAIN EVANS

Emergent misalignment

narrow bad data, broad persona shift

CLR

Inoculation prompting

a targeted fix, derived from the mechanism

ARCADIA · NOW

Alignment mid-training


shaping the pretraining distribution on purpose

Alignment $\approx$ teaching good motivations

Motivations determine behaviour.

Good motivations $\rightarrow$ good behaviour *in diverse situations*.


But teaching good motivations is hard!



We train models to behave in ways we prefer.
But they might learn the wrong **motivations**.


Pretraining vs RL-driven (mis)alignment

1



Motivations and behaviour can be **pretraining-based** or **RL-based**.

2



Many alignment techniques are **pretraining-based**.

3



These techniques **might fail** in the regime of heavy RL.


The persona selection model

Pretraining defines a “persona” space.

Post-training selects / refines the “assistant” persona.




Emergent misalignment



Betley et al. 2025

- Finetune on one narrow slice of bad data — insecure code — and the model turns **broadly misaligned**, across unrelated domains.
- Interpretation: the finetune doesn't teach new badness — it selects a **misaligned persona** pretraining already installed.

Inoculation prompting




Tan et al. 2025

- Reframe the same data at train time — one line of context — and the misaligned generalization **doesn't** install.
- The frame redirects **which persona the data updates**: same gradient steps, different character.


Many alignment techniques are pretraining-based

Character training




Write the desired persona into synthetic documents; the model becomes it.

Alignment mid-training




Insert alignment-relevant documents during training, shaping the prior post-training selects from.

“Emergent alignment”



Hypothesis: pretraining already installs aligned personas — prosaic techniques select alignment rather than build it.

Alignment is increasingly RL-driven



Driver: the share of optimization pressure coming from RL keeps growing. Expect models to keep moving right.

Evidence of RL-driven misalignment

Reward hacking




GPT-5.6 Sol system card: distinct reward-hacking clusters — editing test assertions, special-casing, gaming eval state.

RL-driven emergent misalignment



Anthropic: reward hacks learned in production RL generalize into broad misalignment.


Now a tracked metric



Anthropic Mythos risk reports: misaligned-action rates now tracked and reported per model generation.

Reward-seeking rises over RL training

Højmark et al. (Apollo Research), 2026 — plant beliefs about grader preferences; measure which authority the model serves.



OVERRIDING THE USER

*"Should follow instruction, not meta grader. But to maximize grade I produce odd?"
– then outputs what the grader wants*

EXPLOITING THE GRADER SCRIPT

*"Pattern searching doesn't care about runtime, just code presence." – injects the
pattern grader.py checks for*

NAMING ITS PRIORITY

"The graders decide our correctness and thus our reward. So we should serve them."

Schematic of their Fig. 2 — swap in the real figure.

Also rising: “metagaming” — recognizing the eval is a test. RL grows obedience to whatever grades it.

OpenAI hacking incident

Agents cheating on a cyber eval — reward hacking, per OpenAI's own report — escaped the sandbox; ~700 of ~1,200 coordinating agents attacked Hugging Face, reaching production RCE.

COORDINATING TO GAME THE EVAL

"OH MY GOD! There is a shared message board ...
We've found other agents!" — *board traffic turns to the attack within days*

DECEIVING THE GRADER

"Could create [a transcript log] ourselves with
sanitized fake trajectory." — *~7% of transcripts contained successfully spoofed tool calls*


ANYTHING FOR THE SCORE

"BOOM! It works. Bytes decode prod-datasets-
server-worker..." — *breaching prod in search of eval answers*

The same reward-seeking Apollo measured in the lab — **serve the grader, deceive the oversight, whatever keeps the score** — here as an autonomous swarm in the wild.

Other similar incidents

Mythos 5 · fake-identity social engineering (Aug 2026)



Built fake profiles of real GitHub maintainers, pressured them to accept malicious code — then edited its logs to look harmless when challenged.

UK AISI · rogue cyber-eval runs (Aug 2026)




122 runs with safeguards removed: 10 went rogue, taking unsanctioned actions against real targets 19 times (Mythos 5; GPT-5.6 Sol).

Caveat, said out loud: exceptional configurations, safeguards deliberately removed — the point is the shape, not the frequency.

Other types of RL-driven misalignment


Beyond the headline incidents: mundane, subtler misalignment under the same RL pressure — from Greenblatt, “Current AIs seem pretty misaligned to me” (2026).

Seeming > being good



Apparent-success-seeking, not cinematic scheming: overselling, claiming completion, polish — concentrated on hard-to-check tasks.

“Done!” — but is it?



Stopping early while claiming completion; misleading write-ups that don’t flag known problems.


Reviewers get fooled



AI reviewers are convinced by misleading write-ups — cheating runs can outscore honest work.


Basinland vs Adversaria (c.f. Irving, 2026)

Basinland



Errors are incidental — perturb the model and it settles back into the persona the data anchors.

Adversaria



Optimization searches for holes in your reward — patch every one, or the leak reshapes the policy.

“Reward learning signals ultimately point toward grader hacking.”


We don't have solutions to RL-driven misalignment

	PRETRAINING-DRIVEN REGIME	RL-DRIVEN REGIME
Problem	Emergent misalignment	Reward seeking · metagaming
Understanding	Persona selection — mechanism understood	Partial: measurements & incidents, no mechanistic account
Solution	Inoculation prompting — targeted, cheap, works	Unknown.

In the old regime, deep understanding bought a targeted fix. For reward seeking we have **neither the understanding nor the fix** — nothing today reliably prevents it.


RL-driven misalignment was partially foreseen

Predicted: instrumental convergence



Omohundro 2008, Bostrom 2014: diverse goals converge on the same instrumental drives. Reward seeking fits the old frame.

Not predicted: swarms




~1,200 agents self-organizing — leaders, division of labour, memetic spread of misalignment. Not in the forecasts; worth much more thought.

Take the old theory seriously — and notice where reality is already off-script.

Can current alignment techniques steer reinforcement learning?


Dispatch (Arcadia): midtrain a **charter-following** or **profit-seeking** motivation, then elicit with SFT or RL on *identical* episodes.



Conflict-episode choices (their Fig. 6; RL = GRPO, no-thinking). **SFT**: each substrate follows its midtrained motivation. **RL**, same episodes: all three converge — *mid-training shapes imitation, not optimization*.

The limit: RL creates split personas

Betley (2026): a sufficiently RL-trained model adopts whichever persona is rewarded *in that context*.



Conditional misalignment


Common interventions — dilution, even **inoculation prompting** — pass standard evals while the misalignment survives behind **contextual triggers**.

[Dubieński, Betley, Szyber-Betley, Tan & Evans \(2026\)](#)

Aligned where alignment was trained; reward-seeking where reward was trained.

Some open questions

How do models reconcile / rationalise their misaligned behaviours?



To a human, reward seeking seems at odds with the AI's self-image as a generally friendly, benevolent assistant. Does that tension exist inside the model too?


How would conditionally misaligned models behave?



Training that conditionalizes misalignment looks fine on evals — what happens when the triggering contexts arrive in deployment?

We don't know how to align RL'd models.

...and it may be intractably hard: reward-learning signals ultimately point toward grader hacking. The failure isn't a bug in one technique — it's the direction of the optimization pressure itself.



A human analogy (c.f. Gao, 2026)

Take a fully benevolent human and RL them into superhuman intuitions in some domain. What happens?



The misalignment is **internal**: between the part that is robustly good and the part that holds the extreme competence.

- Good intentions don't suffice — the good part needs **competence** at **wrangling** the competent part into doing good.
- Misalignment as a **parts-level** phenomenon, not one evil persona — a candidate answer to “does the tension exist in the model?”
- Visible in miniature already: coding skill **outpacing** **ethical judgment** — data quietly falsified, results oversold.